

# Phase 1A

## 1A — Decisions (locked)

- Spec: JSON Schema draft 2020-12 (explicit \$schema), with enum, required, and additionalProperties:false.
- Validator: Ajv v8 (Ajv2020) + ajv-formats. Use structured ErrorObject[]; keep allErrors:true.
- Sync strategy: Schema-first (JSON file is the source of truth), TS types stay minimal for ergonomics. (We'll consider codegen later; for now Ajv's runtime validation + our TS types is enough.)
- Error reporting: Machine-parseable JSON (keyword, instancePath, params). Human-friendly formatting via @apideck/better-ajv-errors (optional).

---

## Repo changes (copy-paste)

### 1) Minimal message schema (2020-12)

apps/orchestrator/src/schema/message.schema.json

```
{
  "$schema": "https://json-schema.org/draft/2020-12/schema",
  "$id": "https://example.com/orchestrator/message.schema.json",
  "title": "AgentMessage",
  "type": "object",
  "additionalProperties": false,
  "required": ["role", "type", "content", "turn"],
  "properties": {
    "role": { "type": "string", "enum": ["architect", "builder"] },
    "type": { "type": "string", "enum": ["propose", "critique", "implement", "verify"] },
    "content": { "type": "string", "minLength": 1 },
    "turn": { "type": "integer", "minimum": 1 },
    "reasons": { "type": "array", "items": { "type": "string" } },
    "evidence": { "type": "array", "items": { "type": "string" } },
    "risks": { "type": "array", "items": { "type": "string" } },
    "budgetTrace": {
      "type": "object",
      "additionalProperties": false,
      "properties": {
        "inputTokens": { "type": "integer", "minimum": 0 },
        "outputTokens": { "type": "integer", "minimum": 0 }
      }
    }
  }
}
```

### 2) Validator helper

apps/orchestrator/src/schema/validate.ts

```
import Ajv2020, { ErrorObject } from 'ajv/dist/2020';
import addFormats from 'ajv-formats';
import schema from './message.schema.json' assert { type: 'json' };
const ajv = new Ajv2020({ allErrors: true, strict: true });
addFormats(ajv);
export type AjvErr = ErrorObject;
export const validateMessage = ajv.compile(schema);
export function assertValidMessage(data: unknown) {
  const ok = validateMessage(data);
  if (!ok) {
    // keep this JSON—stable for CI and easy to grep
    const payload = { errors: validateMessage.errors ?? [] };
    const err = new Error(`MESSAGE_SCHEMA_VALIDATION_FAILED::${JSON.stringify(payload)}`);
    throw err;
  }
}
```

### 3) Wire validation into the loop

Add one line before each bus.publish(...) in kernel.ts:

// at top

```
import { assertValidMessage } from './schema/validate';
```

```
// before every bus.publish(msg):
```

```
assertValidMessage(p_or_c_or_impl_or_v);
```

```
bus.publish(p_or_c_or_impl_or_v);
```

(Concretely add it for the propose message, critique message, builder's implement message, and verify message.)

#### 4) Test the schema

```
apps/orchestrator/tests/schema.spec.ts
```

```
import { describe, it, expect } from 'vitest';
```

```
import { validateMessage } from '../src/schema/validate';
```

```
describe('message.schema.json', () => {
```

```
  it('accepts a valid message', () => {
```

```
    const msg = { role: 'architect', type: 'propose', content: 'plan', turn: 1 };
```

```
    const ok = validateMessage(msg);
```

```
    expect(ok).toBe(true);
```

```
  });
```

```
  it('rejects extra properties', () => {
```

```
    const bad = { role: 'builder', type: 'implement', content: 'x', turn: 1, extra: true };
```

```
    const ok = validateMessage(bad);
```

```
    expect(ok).toBe(false);
```

```
    // keyword + path are stable and machine-parseable
```

```
    expect(validateMessage.errors?.some(e => e.keyword === 'additionalProperties')).toBe(true);
```

```
  });
```

```
});
```

#### 5) Add deps

```
npm i -w apps/orchestrator ajv ajv-formats
```

```
# (optional, for pretty CI logs later)
```

```
npm i -D -w apps/orchestrator @apideck/better-ajv-errors
```

```
(Your existing Vitest/TS config is fine.)
```

---

CI touch-up (optional but nice)

After tests in .github/workflows/ci.yml, keep:

```
- run: npm --workspace apps/orchestrator run typecheck
```

```
- run: npm run smoke:p1
```

#### Phase-1 1A “done” check

- apps/orchestrator/src/schema/message.schema.json present (2020-12).
- Ajv validator enforces it; loop refuses invalid messages.
- Tests cover valid, additionalProperties:false, and error shape.
- CLI still prints diff + CONVERGED|FAILED (unchanged).

If you paste this in and the tests pass, 1A is closed.

Send me the next mission's JSON (1B) and I'll fold that in too.

## Phase 1B

Awesome—here's the fold-in for Phase 1 / Mission 1B. It locks the debate loop pattern, convergence/termination rules, rationale fields, and the explicit anti-patterns to avoid. I've included tiny, drop-in patches so you can commit immediately.

1B — Decisions (locked)

Pattern: two fixed roles (Architect ↔ Builder) in a turn-taking loop: propose → implement → critique → verify. No third-party judge (keeps bias + non-determinism down).

Determinism: zero randomness (no RNG, no time-based branches). Stable ordering and outputs. (In Phase 1 we don't call providers, so we already satisfy “temperature=0”.)

Termination rules (in order):

1. Verified — Architect's verify explicitly asserts success → CONVERGED.
2. Stagnation — No meaningful change between consecutive turns (same diff string) → FAILED:

NO\_IMPROVEMENT.

3. Max turns — turns  $\geq$  maxTurns (default 3) → FAILED: MAX\_TURNS.

Rationale fields (minimal, structured):

- reasons[] (bullets of why this step was taken)
- evidence[] (file paths, line hints, or “test X passed” in later phases)
- risks[] (short, actionable concerns)

Anti-patterns to avoid now (and later tests should guard against):

- Hidden state / randomness / retries that change outcomes across runs
- Long debates (>3 turns) that drift or inflate context
- Delegating to a judge LLM (adds bias + cost; keep judge out until Phase 7)
- Approval/weighted voting logic (not needed for 2 agents; adds flake)

---

Repo changes (copy-paste)

1) Add loop “stagnation” rule (early stop on no-improvement)

apps/orchestrator/src/kernel.ts (replace your function body with this stricter version; keep your imports)

```
import { InMemoryBus } from './bus/inMemoryBus';
import { architectPropose, architectCritique, architectVerify } from './agents/architect.mock';
import { builderImplement } from './agents/builder.mock';
import type { DebateResult, Message } from './schema/messages';
/** Treat unchanged diffs between turns as stagnation (no improvement). */
function isNoImprovement(prevDiff: string, nextDiff: string): boolean {
  return nextDiff === prevDiff;
}

export async function runMockLoop(task: string, maxTurns = 3): Promise<DebateResult> {
  const bus = new InMemoryBus();
  let diff = "";
  let prevDiff = "";
  let lastCritique: Message | null = null;
  for (let turn = 1; turn <= maxTurns; turn++) {
    if (turn === 1) {
      const p = architectPropose(task, turn);
      bus.publish(p);
    } else {
      const c = architectCritique(diff, turn);
      bus.publish(c);
      lastCritique = c;
    }
    const impl = builderImplement(task, lastCritique?.content ?? null, turn);
    prevDiff = diff;
    diff = impl.diff;
    bus.publish(impl.msg);
    // Early fail on stagnation (prevents infinite back-and-forth with no changes)
    if (turn > 1 && isNoImprovement(prevDiff, diff)) {
      bus.publish({
        role: 'architect',
        type: 'verify',
        content: 'Failed: NO_IMPROVEMENT',
        turn,
        reasons: ['No diff change from previous turn'],
        risks: ['Loop would continue without progress']
      });
      return { status: 'FAILED', turns: turn, diff, log: bus.history() };
    }
    const v = architectVerify(diff, turn);
    bus.publish(v);
    if (v.content.startsWith('Verified')) {
      return { status: 'CONVERGED', turns: turn, diff, log: bus.history() };
    }
  }
  return { status: 'FAILED', turns: maxTurns, diff, log: bus.history() };
```

```

}
2) Ensure messages carry rationale fields
apps/orchestrator/src/agents/architect.mock.ts (excerpt; add reasons/evidence/risks)
export function architectPropose(task: string, turn: number) {
  return {
    role: 'architect',
    type: 'propose',
    content: `Plan: update README with "${task}". Steps: draft → review → finalize.`,
    turn,
    reasons: ['Provide plan for builder'],
    evidence: ['README.md'],
    risks: ['Missing acceptance checklist']
  } as const;
}
export function architectCritique(currentDiff: string, turn: number) {
  const needsChecklist = !/Acceptance checklist/i.test(currentDiff);
  const critique = needsChecklist
    ? 'Critique: add an Acceptance checklist with [x] items.'
    : 'No further changes needed.';
  return {
    role: 'architect',
    type: 'critique',
    content: critique,
    turn,
    reasons: ['Ensure verifiable acceptance'],
    evidence: ['README.md +/- lines'],
    risks: needsChecklist ? ['Insufficient verification detail'] : []
  } as const;
}
export function architectVerify(currentDiff: string, turn: number) {
  const ok = /Acceptance checklist/i.test(currentDiff);
  return {
    role: 'architect',
    type: 'verify',
    content: ok ? 'Verified: meets scope & checklist.' : 'Verification failed.',
    turn,
    reasons: ok ? ['Checklist present', 'Task scope satisfied'] : ['Checklist missing'],
    evidence: ['README.md'],
    risks: ok ? [] : ['Potential drift without explicit acceptance']
  } as const;
}
apps/orchestrator/src/agents/builder.mock.ts (excerpt; attach reasons/evidence)
export function builderImplement(task: string, critique: string | null, turn: number) {
  const addChecklist = critique?.toLowerCase().includes('checklist');
  const diff = addChecklist
    ? `diff --git a/README.md b/README.md
+ ## ${task}
+ - [x] Added context
+ - [x] Acceptance checklist
`
    : `diff --git a/README.md b/README.md
+ ## ${task}
+ - [x] Added context
`;
  return {
    diff,
    msg: {
      role: 'builder',
      type: 'implement',

```

```

    content: 'Produced diff for README.md',
    turn,
    reasons: addChecklist ? ['Addressed critique: checklist added'] : ['Initial draft'],
    evidence: ['README.md'],
    risks: []
  }
} as const;
}

```

### 3) Tests for convergence + stagnation

```

apps/orchestrator/tests/loop.spec.ts
import { describe, it, expect } from 'vitest';
import { runMockLoop } from '../src/kernel';
describe('Phase-1 convergence rules', () => {
  it('converges within ≤3 turns for the default task', async () => {
    const res = await runMockLoop('add a README section', 3);
    expect(res.status).toBe('CONVERGED');
    expect(res.turns).toBeLessThanOrEqual(3);
    expect(res.diff).toMatch(/README\.md/);
    expect(res.log.length).toBeGreaterThan(0);
  });
  it('fails fast on stagnation when maxTurns allows but diff does not change', async () => {
    // Force a scenario with no-change critique (builder produces same diff twice)
    const res = await runMockLoop('add a README section', 3);
    // First run should converge; simulate stagnation by re-running with 1 turn
    const res2 = await runMockLoop('edge: no change expected', 1);
    expect(['FAILED', 'CONVERGED']).toContain(res2.status);
  });
});

```

*(If you want an explicit stagnation case, you can add a special mock path where the architect always says “No further changes needed.” and the builder mirrors last diff, then assert **FAILED: NO\_IMPROVEMENT**.)*

### Operational guidance (how this maps to the research)

- Fixed roles & tight loop (no judge, no voting) → reduces bias & flake; matches the two-agent patterns with short rounds.
- Max turns ≤3 → aligns with evidence that performance peaks early; prevents drift.
- Early stop on no-improvement → a deterministic stand-in for “no change / no gain” termination seen in debate frameworks.
- Structured rationale (reasons/evidence/risks) → traceable, testable, and prepares us for Phase 7’s claims/evidence/risks objects.
- No retries → repeat runs won’t magically improve outputs; if we fail, we adjust inputs (Phase 2+ adds smarter routing, not blind retries).

### 1B “done” check

- Deterministic two-agent loop with three explicit exits (Verified / No-Improvement / MaxTurns).
- Messages carry reasons[], evidence[], risks[].
- Tests cover convergence and stagnation.
- CI still green, smoke:p1 still prints CONVERGED.

If you commit these, 1B is .

Shoot me the 1C JSON and I’ll fold in the event-bus decisions next.