

Phase 1 plan

Here's a crisp, end-to-end plan for Phase 1 — Two-Agent Loop (All Mocks) that maps straight onto your repo and promotion rules.

Phase 1 — Two-Agent Loop (All Mocks)

1) Objective

Prove a typed, traceable, two-agent debate loop locally (no real model calls). Agents exchange structured messages (propose → implement → critique → verify) via an in-memory bus; the loop converges in ≤ 3 turns and produces a mock diff.

2) Scope & Non-Goals

In scope

- JSON message schema + TS types
- In-memory event bus
- Deterministic loop exits (converged / maxTurns / failure)
- CLI that prints a diff and CONVERGED|FAILED
- Unit tests + CI

Not in scope (Phase 2+)

- Real provider adapters (Anthropic/OpenAI)
- Git PR automation
- RAG, review gates, policy engine

3) Deliverables & Repo Changes

- **Schema**
- apps/orchestrator/src/schema/messages.ts (**TS types**)
- apps/orchestrator/src/schema/message.schema.json (**JSON Schema 2020-12**)
- **Loop & Bus**
- apps/orchestrator/src/bus/inMemoryBus.ts
- apps/orchestrator/src/kernel.ts (**loop logic & exits**)
- apps/orchestrator/src/agents/*.mock.ts (**architect/builder**)
- **CLI**
- apps/orchestrator/src/cli/pocMock.ts (**prints diff + status; proper exit code**)
- **Scripts**
- apps/orchestrator/package.json: "poc:mock", "typecheck"
- **Root** package.json: "smoke:p1": "npm --workspace apps/orchestrator run poc:mock"
- **Tests**
- apps/orchestrator/tests/kernel.spec.ts
- apps/orchestrator/tests/poc.spec.ts (**fail & maxTurns**)
- **Docs**
- docs/architecture/phase-1-plan.md (**this plan**)
- **CI**
- .github/workflows/ci.yml: **add** typecheck + smoke:p1

4) Protocol (message contract)

TS Types (apps/orchestrator/src/schema/messages.ts)

```
export type MessageType = 'propose' | 'critique' | 'implement' | 'verify';
export type Role = 'architect' | 'builder';
export interface Message {
  role: Role;
  type: MessageType;
  content: string;
  turn: number;
  reasons?: string[]; // optional rationale bullets
  evidence?: string[]; // optional file/line refs
  risks?: string[]; // optional risk calls
  budgetTrace?: { inputTokens?: number; outputTokens?: number };
}
export interface DebateResult {
  status: 'CONVERGED' | 'FAILED';
  turns: number;
  diff?: string;
```

```
log: Message[];  
}
```

JSON Schema (apps/orchestrator/src/schema/message.schema.json)

- **Draft 2020-12**
- **Required:** role, type, content, turn
- **Disallow additional properties (keeps mocks tight)**

5) Loop Algorithm (deterministic)

Pseudocode

diff = ""

lastCritique = null

for turn in 1..maxTurns:

 if turn == 1: bus.publish(architect.propose(task))

 else: lastCritique = architect.critique(diff); bus.publish(lastCritique)

 impl = builder.implement(task, lastCritique?.content); diff = impl.diff; bus.publish(impl.msg)

 verify = architect.verify(diff); bus.publish(verify)

 if verify says "Verified": return CONVERGED with diff + log

return FAILED with diff + log

Exit conditions

- **Converged:** verify contains "Verified"
- **Timeout:** turns >= maxTurns → FAILED
- **Budget abort (stub):** if set, force FAILED (placeholder for Phase 2+)

6) In-Memory Bus

- publish(msg) appends to an array
- history() returns a copy (used for result.log)
- **No concurrency; single-thread event order guarantees test determinism**

7) CLI & Scripts

CLI apps/orchestrator/src/cli/pocMock.ts

- **Args:** free-form task ("add a README section" default)
- **Env:** MAX_TURNS (default 3)
- **Prints** diff then CONVERGED|FAILED, **exits** 0/1

Scripts

- apps/orchestrator/package.json
- "poc:mock": "tsx src/cli/pocMock.ts \"\${TASK:-add a README section}\""
- "typecheck": "tsc -p tsconfig.json --noEmit"
- **Root** package.json
- "smoke:p1": "npm --workspace apps/orchestrator run poc:mock"

8) Tests (must pass locally & in CI)

1. **Converge:** default task converges ≤3 turns; diff contains README.md
2. **Forced fail:** maxTurns=1 returns FAILED
3. **Log order:** bus preserves chronological order (architect→builder→verify...)
4. **Schema (optional):** validate a sample Message against message.schema.json
5. **Negative (optional):** malformed message rejected (if you add a validator)

9) CI Changes

Append to the job after npm test:

- run: npm --workspace apps/orchestrator run typecheck

- run: npm run smoke:p1

(Keep existing checks: build-test, CodeQL, Scorecard.)

10) Smoke Test (promotion gate)

From repo root:

npm run smoke:p1 *# prints diff + CONVERGED, exit 0*

npm --workspace apps/orchestrator run typecheck

npm test

Definition of Done (Phase 1):

- poc:mock **converges within ≤ 3 turns** for "add a README section"
 - **Unit tests pass; CI is green** including typecheck and smoke:p1
 - **Message schema (TS + JSON Schema) present; bus + exits implemented**
 - **A 1-page docs/architecture/phase-1-plan.md exists**
-

11) Budget & Telemetry (stubbed)

- **Keep** budgetTrace fields in Message; **populate with zeros (no real calls)**
 - **Ensure the CLI prints only diff + status (no spend), preserving Phase 0E guardrails for later**
-

12) Risks & Mitigations

- **Loop never converges** → add stricter verification rule & maxTurns fail path (already in)
 - **Schema drift** → lock with JSON Schema; add a sample validation test
 - **Flaky tests** → pure functions; no timers; single-thread bus
 - **Scope creep** → no API calls or Git touches in Phase 1
-

13) PR Checklist (use CodeRabbit + Copilot)

- **Schema files added & referenced**
 - **CLI added; poc:mock script wired**
 - **Tests cover converge + fail + ordering**
 - **CI runs typecheck + smoke:p1**
 - docs/architecture/phase-1-plan.md **created**
 - **All checks green; branch policy satisfied**
-

14) "Do now" (three commands)

git checkout -b phase-1/bootstrap

npm run smoke:p1 || true *# ensure script wired; fix until it prints CONVERGED*

npm test

If anything blocks, lean on Copilot for inline fixes, Codex for scaffolds, and open a PR for CodeRabbit to review.

15) Promotion to Phase 2 (readiness)

Promote when:

- **Phase 1 DoD met (above)**
- **Report in docs/architecture/phase-1-plan.md includes:**
- **Protocol summary**
- **Exit rules**
- **Test matrix & CI proof (short excerpts)**

Then we move to Phase 2 — Real Provider Adapters (Claude + OpenAI/GPT).